


```
In [2]: # demonstration of the central limit theorem
from numpy.random import seed
from numpy.random import randint
from numpy import mean
from matplotlib import pyplot
```

```
In [3]: # seed the random number generator
seed(1)


def plot_clt(n):
    # calculate the mean of 50 dice rolls n times
    means = [mean(randint(1, 7, 50)) for _ in range(n)]

    # plot the distribution of sample means
    pyplot.hist(means)
    pyplot.show()
```


```
In [4]: plot_clt(100)
```



```
In [5]: plot_clt(1000)
```



In [6]: `plot_clt(10000)`



```

In [2]: #Effect of Standard deviation on Margin of error
#MOE = z_alpha * sigma/sqrt(n)
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import math
from scipy.stats import norm


#sigma = 10
z_score = 1.96 # 95% CL
n = 64

moe = [] #List of margin of errors
x = range(0,21) # Varying SD from 0 to 20

for sd in x:
    moe.append(z_score * sd/math.sqrt(n))

plt.plot(x, moe)
plt.title('n = 64, CL = 95% and Varying the Standard deviation')
plt.ylabel('Margin of error')
plt.xlabel('Standard deviation')
plt.show()

```



```

In [3]: #Effect of sample size on Margin of error
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import math
from scipy.stats import norm

```

```

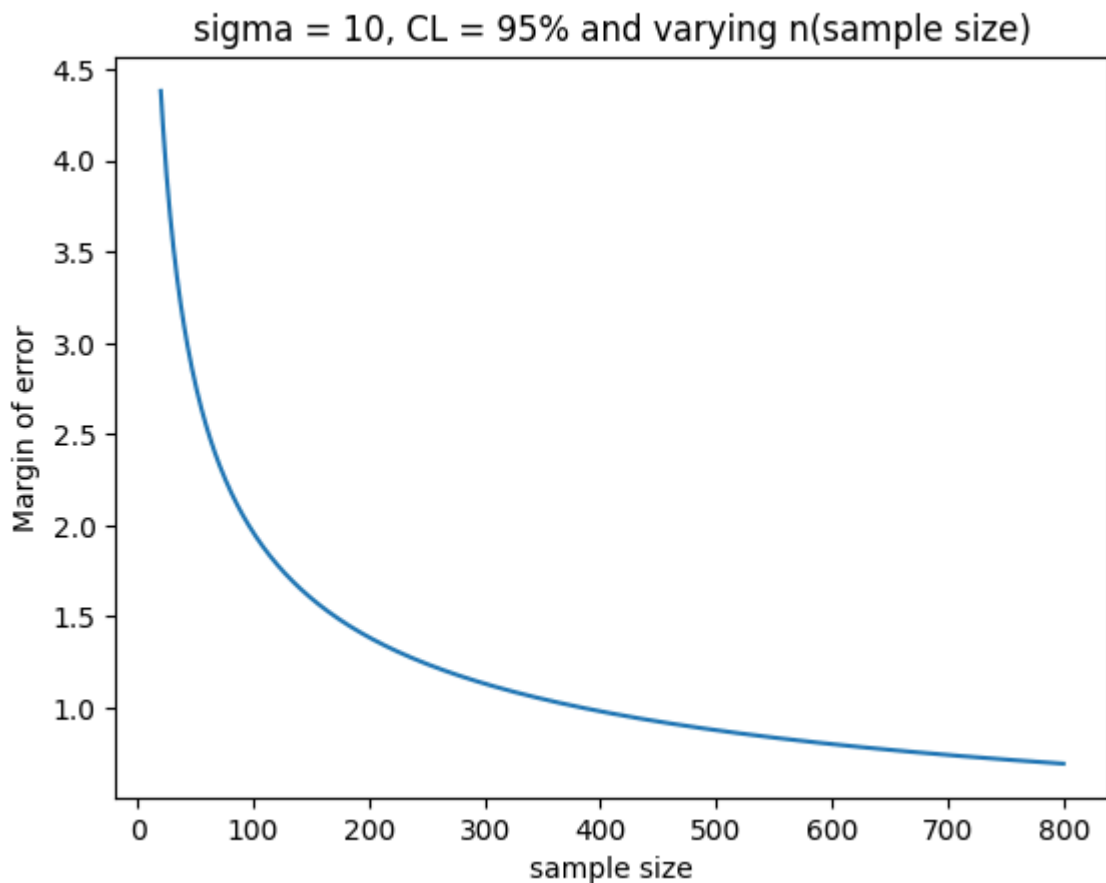
sigma = 10
z_score = 1.96 # 95% CL

moe = []
x = range(20,801) # varying sample size from 20 to 800

for n in x:
    moe.append(1.96 * sigma/math.sqrt(n))

plt.plot(x, moe)
plt.title('sigma = 10, CL = 95% and varying n(sample size)')
plt.ylabel('Margin of error')
plt.xlabel('sample size')
plt.show()

```



```

In [1]: #Effect of Confidence Level on Margin of error
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import math
from scipy.stats import norm

sigma = 10
#z_score = 1.96 # 95% CL
n = 64

def prob(z1,z2):
    return (norm.cdf(z2) - norm.cdf(z1))

z_alphaby2 = np.linspace(norm.ppf(0.10), norm.ppf(0.0005), 100)

```

```

z_alphaby2 = z_alphaby2 * -1

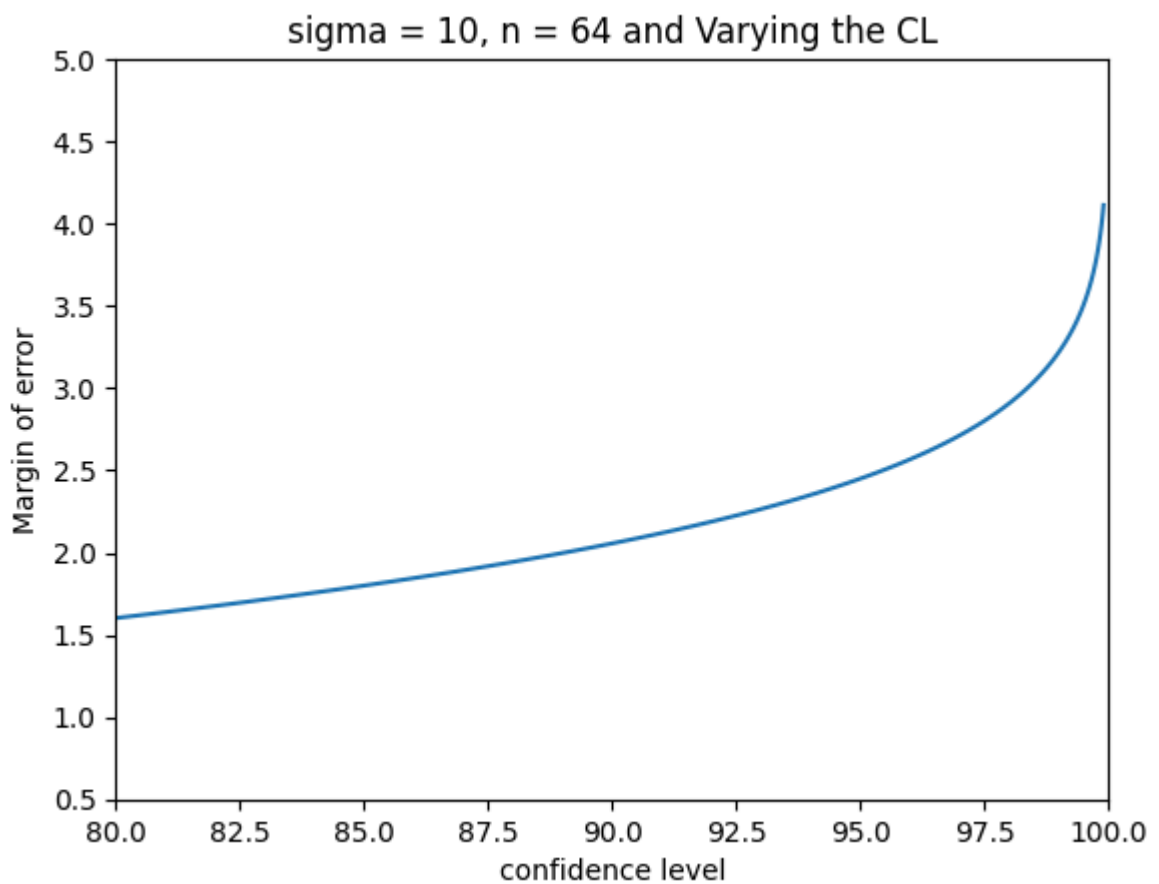
#calculating margin of error
moe = z_alphaby2 * sigma/math.sqrt(n)

#getting the CL
x = []
for i in range(0, len(z_alphaby2)):
    x.append(prob(-1 * z_alphaby2[i], z_alphaby2[i]))

#getting the CL in %
x = np.array(x) * 100

plt.plot(x, moe)
plt.ylabel('Margin of error')
plt.ylim(0.5,5)
plt.xlim(80,100)
plt.title('sigma = 10, n = 64 and Varying the CL')
plt.xlabel('confidence level')
plt.show()

```



In [ ]:

```
In [1]: from scipy.stats import norm
from math import sqrt

def two_sided_hypo(sample_mean, pop_mean, std_dev, sample_size, alpha):
    actual_z = abs(norm.ppf(alpha/2))
    hypo_z = (sample_mean - pop_mean) / (std_dev/sqrt(sample_size))
    print('actual z value :', actual_z)
    print('hypothesis z value :', hypo_z, '\n')
    if hypo_z >= actual_z or hypo_z <= -(actual_z):
        return True
    else:
        return False

alpha = 0.05
sample_mean = 585
pop_mean = 558
sample_size = 100
std_dev = 139

print('H0 :  $\mu =$ ', pop_mean)
print('H1 :  $\mu \neq$ ', pop_mean)
print('alpha value is :', alpha, '\n')

reject = two_sided_hypo(sample_mean, pop_mean, std_dev, sample_size, alpha)
if reject:
    print('Reject NULL hypothesis')
else:
    print('Failed to reject NULL hypothesis')
#variation with different parameters can be shown here
```

H0 :  $\mu =$  558  
H1 :  $\mu \neq$  558  
alpha value is : 0.05

actual z value : 1.9599639845400545  
hypothesis z value : 1.9424460431654675

Failed to reject NULL hypothesis

```
In [2]: #one sided hypothesis test(for smaller than in NULL hypothesis)
def one_sided_hypo(sample_mean, pop_mean, std_dev, sample_size, alpha):
    actual_z = abs(norm.ppf(alpha))
    hypo_z = (sample_mean - pop_mean) / (std_dev/sqrt(sample_size))
    print('actual z value :', actual_z)
    print('hypothesis z value :', hypo_z, '\n')
    if hypo_z >= actual_z:
        return True
    else:
        return False

alpha = 0.05
sample_mean = 108
pop_mean = 100
sample_size = 36
std_dev = 15

print('H0 :  $\mu \leq$ ', pop_mean)
```

```
print('H1 :  $\mu >$ ', pop_mean)
print('alpha value is :', alpha, '\n')

reject = one_sided_hypo(sample_mean, pop_mean, std_dev, sample_size, alpha)
if reject:
    print('Reject NULL hypothesis')
else:
    print('Failed to reject NULL hypothesis')
#variation with different parameters can be shown here
```

H0 :  $\mu \leq 100$

H1 :  $\mu > 100$

alpha value is : 0.05

actual z value : 1.6448536269514729

hypothesis z value : 3.2

Reject NULL hypothesis

```
In [2]: import csv
import pandas as pd
from random import sample

df = pd.read_csv('train.csv')
```

```
In [3]: #simple random sampling
no_of_elements = 10
random_index = sample(range(df.shape[0]), no_of_elements)
print(random_index)
print(df.iloc[random_index])
```

[247, 73, 86, 105, 480, 510, 193, 51, 307, 548]

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
247	LP001819	Male	Yes	1	Not Graduate	No	
73	LP001250	Male	Yes	3+	Not Graduate	No	
86	LP001280	Male	Yes	2	Not Graduate	No	
105	LP001367	Male	Yes	1	Graduate	No	
480	LP002534	Female	No	0	Not Graduate	No	
510	LP002637	Male	No	0	Not Graduate	No	
193	LP001658	Male	No	0	Graduate	No	
51	LP001157	Female	No	0	Graduate	No	
307	LP001994	Female	No	0	Graduate	No	
548	LP002776	Female	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
247	6608	0.0	137.0	180.0	
73	4755	0.0	95.0	NaN	
86	3333	2000.0	99.0	360.0	
105	3052	1030.0	100.0	360.0	
480	4350	0.0	154.0	360.0	
510	3598	1287.0	100.0	360.0	
193	3858	0.0	76.0	360.0	
51	3086	0.0	120.0	360.0	
307	2400	1863.0	104.0	360.0	
548	5000	0.0	103.0	360.0	

	Credit_History	Property_Area	Loan_Status
247	1.0	Urban	Y
73	0.0	Semiurban	N
86	NaN	Semiurban	Y
105	1.0	Urban	Y
480	1.0	Rural	Y
510	1.0	Rural	N
193	1.0	Semiurban	Y
51	1.0	Semiurban	Y
307	0.0	Urban	N
548	0.0	Semiurban	N

```
In [4]: #systematic sampling
Kth = 100
index = [i for i in range(df.shape[0]) if i%Kth==0]
print(df.iloc[index])
```



	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
100	LP001345	Male	Yes	2	Not Graduate	No	
200	LP001674	Male	Yes	1	Not Graduate	No	
300	LP001964	Male	Yes	0	Not Graduate	No	
400	LP002288	Male	Yes	2	Not Graduate	No	
500	LP002603	Female	No	0	Graduate	No	
600	LP002949	Female	No	3+	Graduate	NaN	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
100	4288	3263.0	133.0	180.0	
200	2600	2500.0	90.0	360.0	
300	1800	2934.0	93.0	360.0	
400	2889	0.0	45.0	180.0	
500	645	3683.0	113.0	480.0	
600	416	41667.0	350.0	180.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
100	1.0	Urban	Y
200	1.0	Semiurban	Y
300	0.0	Urban	N
400	0.0	Urban	N
500	1.0	Rural	Y
600	NaN	Urban	N

```
In [5]: #stratified sampling
#stratas formed based on Education
no_of_elements = 4 #number of elements in each strata
unique = list(set(df['Education']))
print('stratas :', unique, '\n')
index_set = [sample(list(df.index[df['Education']==i]), no_of_elements) for i in unique]
index = [j for i in index_set for j in i]
print(df.iloc[index])
```

stratas : ['Not Graduate', 'Graduate']

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
401	LP002296	Male	No	0	Not Graduate	No	
279	LP001908	Female	Yes	0	Not Graduate	No	
66	LP001228	Male	No	0	Not Graduate	No	
200	LP001674	Male	Yes	1	Not Graduate	No	
533	LP002729	Male	No	1	Graduate	No	
132	LP001478	Male	No	0	Graduate	No	
383	LP002234	Male	No	0	Graduate	Yes	
437	LP002401	Male	Yes	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
401	2755	0.0	65.0	300.0	
279	4100	0.0	124.0	360.0	
66	3200	2254.0	126.0	180.0	
200	2600	2500.0	90.0	360.0	
533	11250	0.0	196.0	360.0	
132	2718	0.0	70.0	360.0	
383	7167	0.0	128.0	360.0	
437	2213	1125.0	NaN	360.0	

	Credit_History	Property_Area	Loan_Status
401	1.0	Rural	N
279	NaN	Rural	Y
66	0.0	Urban	N
200	1.0	Semiurban	Y
533	NaN	Semiurban	N
132	1.0	Semiurban	Y
383	1.0	Urban	Y
437	1.0	Urban	Y

```
In [6]: #cluster sampling
#clusters formed based on number of Dependents
no_of_clusters = 5
unique = list(set(df['Dependents']))
smp = sample(unique, no_of_clusters)
print("clusters :", smp, "selected out of :", unique, '\n')
index_set = [list(df.index[df['Dependents']==i]) for i in smp]
index = [j for i in index_set for j in i]
print(df.iloc[index])
```

clusters : [nan, '0', '1', '3+', '2'] selected out of : ['1', nan, '2', '0', '3+', '']

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
6	LP001013	Male	Yes	0	Not Graduate	No	
..	...	...	...	...	...	...	
591	LP002931	Male	Yes	2	Graduate	Yes	
596	LP002941	Male	Yes	2	Not Graduate	Yes	
599	LP002948	Male	Yes	2	Graduate	No	
607	LP002964	Male	Yes	2	Not Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
6	2333	1516.0	95.0	360.0	
..	...	...	...	...	
591	6000	0.0	205.0	240.0	
596	6383	1000.0	187.0	360.0	
599	5780	0.0	192.0	360.0	
607	3987	1411.0	157.0	360.0	
612	7583	0.0	187.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
6	1.0	Urban	Y
..	...	...	...
591	1.0	Semiurban	N
596	1.0	Rural	N
599	1.0	Urban	Y
607	1.0	Rural	Y
612	1.0	Urban	Y

[599 rows x 13 columns]


In [ ]:



```
In [2]: import numpy as np
import matplotlib.pyplot as plt
from random import sample
from statistics import mean

def plot(arr, N, n_samples):
    x = []
    for i in range(1, n_samples):
        #to find N samples from the arr
        smp = sample(arr, N)
        mu = mean(smp)
        x.append(mu)
    plt.hist(x)
    plt.show()

#example data(population)
arr = [i for i in range(1000)]

#variations
plot(arr, 5, 50)
plot(arr, 20, 500)
plot(arr, 50, 500)
#so as number of samples(n_sample) increases the distribution becomes normal
#so as sample size increases the flatness of the distribution decreases
```





In [ ]:


```
In [7]: %matplotlib inline
import matplotlib.pyplot as plt
from scipy.stats import t, norm
import numpy as np
import pandas as pd

x = np.arange(-3.8,4,1/20)    #a random population

for i in [1, 2,10, 29]:
    #plotting all the t-dist curves(pdf gives probab desnity func)
    plt.plot(x, t.pdf(x, i),'--',label=i)

#plotting a regular normal curve
plt.plot(x, norm.pdf(x), label='Std normal curve')
plt.legend(loc = 'upper right')
plt.show()

print("1-cdf gives :", 1-t.cdf(1.59, 2))
print('same as      :', t.sf(1.59, 2))
print(1-norm.cdf(2), norm.sf(2))
```




```
1-cdf gives : 0.12639805893063705
same as      : 0.12639805893063707
0.02275013194817921 0.0227501319481792
```

```
In [8]: def t_table(n, alpha):
    s = t.ppf(alpha/2, n - 1)
    plt.figure(figsize=(8,4))
    plt.plot(x, t.pdf(x, n - 1), color= 'red',label= n - 1)
    #calculating the area under the graphs to be filled
    section1 = np.arange(-5, s, 1/20.)
    section2 = np.arange(-s, 5, 1/20.)
    #fill those above selected areas
    plt.fill_between(section1, t.pdf(section1, n - 1), color='blue')
```

```
plt.fill_between(section2, t.pdf(section2, n - 1), color='blue')
plt.xticks(np.arange(-5,5,0.5), rotation = 45)
plt.legend(loc = 'upper right')
plt.show()


#t_table(sample_size, alpha)
t_table(12, 0.1)
```



```
In [9]: x = np.arange(-7, 8, 1/20)
def ci(t_score, n):
    plt.figure(figsize=(8,4))
    #gives the whole area under the graph
    area = t.cdf(t_score, n - 1) - t.cdf(-t_score, n - 1)
    print('Confidence Level', area * 100)
    plt.plot(x, t.pdf(x, n - 1), color='red',label= n - 1)
    #to fill from -t end to +t end
    section = np.arange(-t_score, t_score, 1/20.)
    plt.fill_between(section, t.pdf(section, n - 1))
    plt.xticks(np.arange(-6,7,0.5), rotation = 45)
    plt.legend(loc = 'upper right')
    plt.show()

ci(5.841, 4)
```

Confidence Level 99.00004355246759





```
In [1]: from scipy.stats import chi2_contingency # defining the table
data = [[207, 282, 241], [234, 242, 232]]
stat, p, dof, expected = chi2_contingency(data) # interpret p-value
alpha = 0.05
print("p value is " + str(p))
if p <= alpha:
    print('Dependent (reject H0)')
else:
    print('Independent (H0 holds true)')
```

p value is 0.10319714047309392  
Independent (H0 holds true)

```
In [2]: import numpy as np
from scipy.stats import chi2

# Observed frequencies
observed = np.array([115, 47, 41, 101, 200, 96])

# Expected frequencies (assuming a fair die)
expected = np.array([100, 100, 100, 100, 100, 100])

# Calculate chi-square statistic
chi2_stat = np.sum((observed - expected)**2 / expected)

# Degrees of freedom (number of categories - 1)
df = len(observed) - 1

# Critical value for 10% significance level
critical_value = chi2.ppf(0.90, df)

# p-value
p_value = 1 - chi2.cdf(chi2_stat, df)

# Output results
print(f"Chi-squared Statistic: {chi2_stat}")
print(f"Critical Value at 10% significance level: {critical_value}")
print(f"p-value: {p_value}")

# Conclusion
if chi2_stat < critical_value:
    print("Fail to reject the null hypothesis: The die is unbiased.")
else:
    print("Reject the null hypothesis: The die is biased.")
```

Chi-squared Statistic: 165.32000000000002  
Critical Value at 10% significance level: 9.236356899781123  
p-value: 0.0  
Reject the null hypothesis: The die is biased.

```
In [3]: import numpy as np
import pandas as pd
from scipy.stats import chi2_contingency

# Define the observed data
data = np.array([
    [10, 102, 8], # Machine 1
    [34, 161, 5], # Machine 2
    [12, 79, 9], # Machine 3
```

```

    [10, 60, 10]    # Machine 4
])

# Create a DataFrame for better visualization (optional)
df = pd.DataFrame(data, columns=['Too Thin', 'OK', 'Too Thick'],
                  index=['Machine 1', 'Machine 2', 'Machine 3', 'Machine 4'])

print("Observed Data:\n", df)

# Perform the Chi-Square test
chi2_stat, p_value, dof, expected = chi2_contingency(data)

# Display results
print("\nChi-Square Statistic:", chi2_stat)
print("P-Value:", p_value)
print("Degrees of Freedom:", dof)
print("Expected Frequencies:\n", expected)

# Determine if the result is significant
alpha = 0.05
if chi2_stat > chi2.ppf(1 - alpha, dof):
    print("Reject the null hypothesis: There is a significant difference.")
else:
    print("Fail to reject the null hypothesis: No significant difference.")

```

Observed Data:

	Too Thin	OK	Too Thick
Machine 1	10	102	8
Machine 2	34	161	5
Machine 3	12	79	9
Machine 4	10	60	10

Chi-Square Statistic: 15.584353328056686

P-Value: 0.01616760116149423

Degrees of Freedom: 6

Expected Frequencies:

```

[[ 15.84  96.48   7.68]
 [ 26.4  160.8  12.8 ]
 [ 13.2   80.4   6.4 ]
 [ 10.56  64.32   5.12]]

```

Reject the null hypothesis: There is a significant difference.

```

In [4]: import numpy as np
import pandas as pd
from scipy.stats import chi2_contingency
import matplotlib.pyplot as plt

# Create a contingency table
data = np.array([[150, 30],    # Vaccinated
                 [80, 40]])   # Not Vaccinated

# Display the contingency table as a DataFrame for clarity
contingency_table = pd.DataFrame(data,
                                columns=['Recovered', 'Not Recovered'],
                                index=['Vaccinated', 'Not Vaccinated'])
print("Contingency Table:\n", contingency_table)

# Perform the Chi-Square test
chi2_stat, p_value, dof, expected = chi2_contingency(data)

```

```

# Display results
print("\nChi-Square Statistic:", chi2_stat)
print("P-Value:", p_value)
print("Degrees of Freedom:", dof)
print("Expected Frequencies:\n", expected)

# Determine significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant association between")
else:
    print("Fail to reject the null hypothesis: No significant association between")

# Optional: Plotting the contingency table
plt.figure(figsize=(8, 5))
plt.title("Vaccination vs Recovery Status")
plt.bar(['Vaccinated', 'Not Vaccinated'], [150, 80], label='Recovered', color='lightblue')
plt.bar(['Vaccinated', 'Not Vaccinated'], [30, 40], label='Not Recovered', color='lightcoral')
plt.ylabel('Number of Patients')
plt.legend()
plt.grid(axis='y')
plt.show()

```

Contingency Table:

	Recovered	Not Recovered
Vaccinated	150	30
Not Vaccinated	80	40

Chi-Square Statistic: 10.267857142857142

P-Value: 0.0013536793727780064


Degrees of Freedom: 1

Expected Frequencies:

[[138. 42.]

[ 92. 28.]]

Reject the null hypothesis: There is a significant association between vaccination and recovery.



```
In [5]: import numpy as np
import pandas as pd
from scipy.stats import chi2_contingency
import matplotlib.pyplot as plt

# Create a contingency table
data = np.array([[30, 10], # Male
                 [20, 30]]) # Female

# Display the contingency table as a DataFrame for clarity
contingency_table = pd.DataFrame(data,
                                columns=['Purchased', 'Not Purchased'],
                                index=['Male', 'Female'])
print("Contingency Table:\n", contingency_table)

# Perform the Chi-Square test
chi2_stat, p_value, dof, expected = chi2_contingency(data)

# Display results
print("\nChi-Square Statistic:", chi2_stat)
print("P-Value:", p_value)
print("Degrees of Freedom:", dof)
print("Expected Frequencies:\n", expected)

# Determine significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant association between")
else:
    print("Fail to reject the null hypothesis: No significant association between")

# Optional: Plotting the contingency table
plt.figure(figsize=(8, 5))
plt.title("Gender vs Product Purchase Preference")
plt.bar(['Male', 'Female'], [30, 20], label='Purchased', color='lightblue')
plt.bar(['Male', 'Female'], [10, 30], label='Not Purchased', color='salmon', bot
plt.ylabel('Number of Individuals')
plt.legend()
plt.grid(axis='y')
plt.show()
```

Contingency Table:

	Purchased	Not Purchased
Male	30	10
Female	20	30

Chi-Square Statistic: 9.6530625


P-Value: 0.001890361677058677

Degrees of Freedom: 1

Expected Frequencies:

```
[[22.22222222 17.77777778]
 [27.77777778 22.22222222]]
```

Reject the null hypothesis: There is a significant association between gender and product preference.



# Question

Examine the correlation between patients' age and blood pressure levels. The aim is to determine if there is a significant relationship between increasing age and higher blood pressure. Use Pearson correlation to quantify the strength and direction of the relationship

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pearsonr, spearmanr
# Sample data (replace this with actual data)
data = {
    'Age': [25, 30, 35, 40, 45, 50, 55, 60, 65, 70],

    # Uncomment one of the following for different types of correlation
    # 'BloodPressure': [120, 125, 130, 135, 140, 145, 150, 155, 160, 165], # Strong
    'BloodPressure': [120, 125, 130, 135, 120, 145, 100, 155, 160, 165], # Moderate
    # 'BloodPressure': [120, 125, 130, 135, 120, 145, 100, 155, 100, 165], # Weak
    # 'BloodPressure': [150, 140, 135, 130, 125, 120, 110, 100, 95, 90] # Strong
}


# Create a DataFrame
df = pd.DataFrame(data)
print("Data:\n", df)
# Plot the data to visualize the relationship
plt.figure(figsize=(8,6))
sns.scatterplot(x='Age', y='BloodPressure', data=df)
plt.title('Scatterplot of Age vs Blood Pressure')
plt.xlabel('Age')
plt.ylabel('Blood Pressure')
plt.grid(True)
plt.show()
pearson_corr, pearson_p = pearsonr(df['Age'], df['BloodPressure']) # Calculate Pearson Correlation
print(f"Pearson Correlation Coefficient: {pearson_corr:.3f}, p-value: {pearson_p:.3f}")

# Interpretation
if pearson_corr > 0:
    if pearson_corr <= 0.5:
        print("Weak positive correlation.")
    elif 0.5 < pearson_corr < 0.8:
        print("Moderate positive correlation.")
    elif pearson_corr >= 0.8:
        print("Strong positive correlation.")
elif pearson_corr < 0:
    if pearson_corr >= -0.5:
        print("Weak negative correlation.")
    elif -0.8 < pearson_corr < -0.5:
        print("Moderate negative correlation.")
    elif pearson_corr <= -0.8:
        print("Strong negative correlation.")
```

```
else:
    print("No correlation.")
```

Data:

	Age	BloodPressure
0	25	120
1	30	125
2	35	130
3	40	135
4	45	120
5	50	145
6	55	100
7	60	155
8	65	160
9	70	165



Pearson Correlation Coefficient: 0.619, p-value: 0.05647

Moderate positive correlation.

```
In [3]: # Import necessary libraries
import pandas as pd
import numpy as np

# Sample data for correlation (strong positive correlation)
data = {
    'Age': [25, 30, 35, 40, 45, 50, 55, 60, 65, 70],
    # Uncomment one of the following for different types of correlation
    # 'BloodPressure': [120, 125, 130, 135, 140, 145, 150, 155, 160, 165], # Stro
    'BloodPressure': [120, 125, 130, 135, 120, 145, 100, 155, 160, 165], # Moder
    # 'BloodPressure': [120, 125, 130, 135, 120, 145, 100, 155, 100, 165], # Weak
    # 'BloodPressure': [150, 140, 135, 130, 125, 120, 110, 100, 95, 90] # Strong
}

# Create a DataFrame
```

```
df = pd.DataFrame(data)

# Compute means
mean_x = np.mean(df['Age'])
mean_y = np.mean(df['BloodPressure'])

# Pearson correlation computation
numerator = np.sum((df['Age'] - mean_x) * (df['BloodPressure'] - mean_y))
denominator_x = np.sqrt(np.sum((df['Age'] - mean_x) ** 2))
denominator_y = np.sqrt(np.sum((df['BloodPressure'] - mean_y) ** 2))
pearson_corr_manual = numerator / (denominator_x * denominator_y)

print(f"Pearson Correlation Coefficient (Manual Calculation): {pearson_corr_manu

# Interpretation
if pearson_corr > 0:
    if pearson_corr <= 0.5:
        print("Weak positive correlation.")
    elif 0.5 < pearson_corr < 0.8:
        print("Moderate positive correlation.")
    elif pearson_corr >= 0.8:
        print("Strong positive correlation.")
elif pearson_corr < 0:
    if pearson_corr >= -0.5:
        print("Weak negative correlation.")
    elif -0.8 < pearson_corr < -0.5:
        print("Moderate negative correlation.")
    elif pearson_corr <= -0.8:
        print("Strong negative correlation.")
else:
    print("No correlation.")
```

Pearson Correlation Coefficient (Manual Calculation): 0.619

Moderate positive correlation.



```

In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from math import sqrt
from scipy.stats import norm
import random

population = np.arange(1, 10**4) #random population
pop_mean = np.mean(population)


def sampling(sample_size, no_of_samples):
    sample_means = []
    intervals = []
    count = 0
    for i in range(no_of_samples):
        #a sample of size sample_size will be taken
        sample = random.sample(list(population), sample_size)
        #mean of the samples appended to sample_means
        sample_means.append(np.mean(sample))
        #ci contains lower and upper bound of interval with 0.95 confidence
        ci = norm.interval(0.95, np.mean(sample),
                           np.std(sample, ddof=1)/sqrt(sample_size))
        intervals.append(ci)
        #upcount only if pop_mean lies in confidence interval
        if pop_mean >= ci[0] and pop_mean <= ci[1]:
            count = count + 1

    print('Proportion of CIs covering Pop mean', count/no_of_samples)
    plt.figure(figsize=(15,5))
    #print the horizontal line which is pop_mean
    plt.hlines(y = pop_mean, xmin = 0, xmax = 100, color = 'r')
    #print the sample lines with their means indicated as 'o'
    plt.errorbar(np.arange(0.1, 100, 1), sample_means, fmt = 'o', yerr = [(upp -
    plt.show()

    #pass sample_size, no_of_samples
    sampling(1000, 100)

```

Proportion of CIs covering Pop mean 0.97



```

In [2]: #CI for population where 85% of the people say YES to a certain question
import numpy as np
import matplotlib.pyplot as plt
from random import sample
import scipy.stats as st

```

```

import math

#parameters....population, required_CI, sample_size, no_of_samples
def CI(pop, ci, samp_size, no_of_samples):
    print("\nfor ci of", ci, "sample_size", samp_size)
    pop_mean = np.mean(pop)
    print('actual mean :', pop_mean)

    #calculation of same using CI
    samp_means = []      #mean of all the samples
    for i in range(no_of_samples):
        samp_means.append(np.mean(sample(pop, samp_size)))


    #calculation of interval
    print('mean of samples :', np.mean(samp_means))
    pop_stdev = np.std(samp_means) / math.sqrt(samp_size)
    z = st.norm.ppf(ci)
    print("confidence interval :", pop_mean, "+-", z*pop_stdev)
    plt.hist(samp_means)
    plt.show()

pop = sample(range(1, 2*10**5), 10**4) #random population generation


```

In [3]: *#varying no\_of\_samples*  
 CI(pop, 0.85, 1000, 200)  
 CI(pop, 0.85, 1000, 500)  
 CI(pop, 0.85, 1000, 1000)  
*#shape of the curve becomes normal as the no of samples increases(samp\_mean bett*


for ci of 0.85 sample\_size 1000  
 actual mean : 99976.1885  
 mean of samples : 100178.85078  
 confidence interval : 99976.1885 +- 60.56608083307446



for ci of 0.85 sample\_size 1000  
 actual mean : 99976.1885  
 mean of samples : 99908.319476  
 confidence interval : 99976.1885 +- 58.12247581667002




for ci of 0.85 sample\_size 1000  
 actual mean : 99976.1885  
 mean of samples : 99984.57400600001  
 confidence interval : 99976.1885 +- 56.60941099970361




In [4]: `#varying sample size`  
`CI(pop, 0.85, 100, 500)`

```
CI(pop, 0.85, 500, 500)
CI(pop, 0.85, 1000, 500)
#reduction in the size of interval as sample_size increases(better approx of pop
```

```
for ci of 0.85 sample_size 100
actual mean : 99976.1885
mean of samples : 99723.44803999999
confidence interval : 99976.1885 +- 591.4420803012979
```



```
for ci of 0.85 sample_size 500
actual mean : 99976.1885
mean of samples : 99960.261576
confidence interval : 99976.1885 +- 123.1700022730698
```




for ci of 0.85 sample\_size 1000

actual mean : 99976.1885

mean of samples : 99926.154728

confidence interval : 99976.1885  $\pm$  54.621342055094026



```

In [1]: import numpy as np
import matplotlib.pyplot as plt
from scipy import stats

# Given data: Weight (lb) (x) and Length (in.) (y)
weight = np.array([
    0.0, 0.2, 0.4, 0.6, 0.8, 1.0,
    1.2, 1.4, 1.6, 1.8,
    2.0, 2.2, 2.4, 2.6, 2.8,
    3.0, 3.2, 3.4, 3.6, 3.8
])

length = np.array([
    5.06, 5.01, 5.12, 5.13, 5.14, 5.16,
    5.25, 5.19, 5.24, 5.46,
    5.40, 5.57, 5.47, 5.53, 5.61,
    5.59, 5.61, 5.75, 5.68, 5.80
])

# Perform Linear regression
slope, intercept, r_value, p_value, std_err = stats.linregress(weight, length)

# Print regression results
print(f"Intercept: {intercept:.4f}")
print(f"Slope: {slope:.4f}")
print(f"R-squared: {r_value**2:.4f}")

# Generate predicted values
length_predicted = intercept + slope * weight

# Plot the results
plt.figure(figsize=(10, 6))
plt.scatter(weight, length, color='blue', label='Observed Data')
plt.plot(weight, length_predicted, color='red', label='Fitted Line')
plt.title('Linear Regression: Weight vs. Length')
plt.xlabel('Weight (lb)')
plt.ylabel('Length (in.)')
plt.legend()
plt.grid()
plt.show()

```

Intercept: 4.9997

Slope: 0.2046

R-squared: 0.9493

